


# Report CA FP2

0410110 林容安 0410137 劉家麟

## A. Describe our implementation algorithm and explain our results


### 1. Sparse CNN



Sparse CNN is the transformation of CNN, which contains few connection after pruning just like above.

As we have done CNN in Final Project Part1, this time we'd like to use sparse CNN, whose filter and inNeu are composed of many zero elements so that we don't need to take care of those connection.

### 2. Sparse Matrix Implementation: COO Format



We knew that a sparse matrix stored in a common way is really a waste of space, and sparse format solves this problem.

From the picture, we can see that the matrix originally costs 30 sizeof(int) to store. After implementing the COO format, the size decreased to 18. In Final Project Part 2, our matrix has a sparsity of approximately 75 percent, so COO format come in handy.

## B. Discuss what kind of optimization you did (it is better or worse?)

As we talked in part A, we use the CUDA architecture to divide convolution for loops into several blocks and threads. We had cut them in different ways.

The original convolution is about this. There are 6 for loops representing [Filter Number](#), [Frame Size](#), [Frame Depth](#) and [Filter Size](#), separately.

```
for(fn = 0; fn < FILTNUM; fn++){
    for(fmy = 0; fmy < FMSIZE; fmy += STRIDE){
        for(fmx = 0; fmX < FMSIZE; fmX += STRIDE){
            sum = 0;
            for(sli = 0; sli < FMDEPTH; sli++){
                for(y = 0; y < FILTSIZE; y++){
                    for(x = 0; x < FILTSIZE; x++){
                        //do convolution
                    }
                }
            }
            //do ReLU
        }
    }
}
```

There are two cases.

### 1. The one we did in FP1 (not considering sparsity)

In this case, we break the whole convCPU into 2 parts: Convolution & ReLU and MaxPooling. Both have 1-D blocks, 2-D threads.

```
dim3 numBlocks(FILTNUM); //128
dim3 threadsPerBlock(FMSIZE,FMSIZE); //27*27

int bx = blockIdx.x; //FILTNUM 128
int tx = threadIdx.x; //FMSIZE 27 x(col)
int ty = threadIdx.y; //FMSIZE 27 y(row)

for (sli = 0; sli < FMDEPTH; sli++){
    for(y = 0; y < FILTSIZE; y++){ // FILTSIZE 5 y(row)
        for(x = 0; x < FILTSIZE; x++){
            // do convolution
        }
    }
}
```

MaxPooling is below.

```

dim3 P_numBlocks(FILTNUM); //128
dim3 P_threadsPerBlock(FMSIZE/3,FMSIZE/3); //9*9

int bx = blockIdx.x; //FILTNUM 128
int fmx = threadIdx.x; //FMSIZE/3 9 x(col)
int fmy = threadIdx.y; //FMSIZE/3 9 y(row)

```

*The result of case1:*

```

[ca57@hsinchu CA2017FP-Part1]$ nvprof ./CNNConvLayer
==1021== NVPROF is profiling process 1021, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1654.07ms
GPU time for executing a typical convolutional layer = 303.853ms
Congratulations! You pass the check.
Speedup: 5.44366

```

Besides, we found out that setting device first by adding `cudaSetDevice(2);` at the beginning of the main code can also obtain increased speedup from 3 to 5.

w/o `cudaSetDevice(2);`

```

[ca57@hsinchu CA2017FP-Part1]$ nvprof ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1624.5ms
==499== NVPROF is profiling process 499, command: ./CNNConvLayer
GPU time for executing a typical convolutional layer = 492.96ms
Congratulations! You pass the check.
Speedup: 3.29541

```

w/ `cudaSetDevice(2);`

```

[ca57@hsinchu CA2017FP-Part1]$ nvprof ./CNNConvLayer
==1021== NVPROF is profiling process 1021, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1654.07ms
GPU time for executing a typical convolutional layer = 303.853ms
Congratulations! You pass the check.
Speedup: 5.44366

```

## 2. Convolution & ReLU with sparse CNN (maxpooling unmodified)

Since the NNZ of each filter of each channel varies, we could only let the GPU compute a little portion ( $FMSIZE \times FMSIZE$ ) simultaneously. Utilizing the characteristics of sparse matrix, in the sparse CNN, we only have to consider the data which is not zero, so the convolution becomes pretty simple.

Though we parallelize smaller portion than we do in case 1, by implementing sparse CNN, we still obtained a speed up 8.49x, which is a better result compared to case 1(5.44x).

**convolutioning with 1-D blocks, 2-D threads\_\_\_\_\_**

```

dim3 numBlocks(1);
dim3 threadPerBlock(FMSIZE,FMSIZE);

```

```

int tx = threadIdx.x; // FM_SIZE 27 x(col)
int ty = threadIdx.y; // FM_SIZE 27 y(row)

for(int fn = 0 ; fn < FILTNUM ; fn ++){
    sum = 0;
    for(int sli = 0; sli < FMDEPTH; sli++){
        for(int idx=0; idx<FiltCooNNZ[fn*FMDEPTH+sli];idx++){
            CooIdx = tmp + idx;
            ifmx = tx + FiltCooCol[CooIdx]; //col
            ifmy = ty + FiltCooRow[CooIdx]; //row
            inNeuIdx = sli*FM_SIZE*FM_SIZE+ifmy*FM_SIZE + ifmx;

            sum += FiltCooData[CooIdx] * InNeu[inNeuIdx];
            __syncthreads();
        }
        tmp = FiltCooNNZ[fn*FMDEPTH + sli] + tmp;
    }

    outNeuIdx = fn * FM_SIZE * FM_SIZE + ty*FM_SIZE + tx;
    if(sum <= 0)
        outNeural[outNeuIdx] = 0;
    else
        outNeural[outNeuIdx] = sum;
}

```

To implement Sparse CNN, we also modified the input neuron matrix as below.

```

for(int i = 0 ; i < FMDEPTH ; i++){
    for(int j = 0 ; j < (FM_SIZE) ; j++){
        for(int k = 0 ; k < (FM_SIZE) ; k++){
            inNeuIdx = i*(FM_SIZE)*(FM_SIZE) + j*(FM_SIZE) +
k;

            inNeuIdy=i*(FM_SIZE)*(FM_SIZE)+(j-5/2)*FM_SIZE+(k-
5/2);

            if(j>=2 && j<=28 &&k>=2 &&
k<=28)tmp=inNeu[inNeuIdy];
            else tmp = 0 ;
            inGNeu[inNeuIdx] = tmp;

        }
    }
}

```

Where FM\_SIZE is the frame size after modification, and the value is 29.

***The result of case 2:***

```

make rca56@Taipei:~/Sparse_CNN_C00$ make run
./CNNConvLayer
CPU time for executing a typical convolutional layer = 1334.05ms
GPU time for executing a typical convolutional layer = 159.195ms
Congratulations! You pass the check.
Speedup: 8.37995

```

Similarly, we tried adding `cudaSetDevice(2);` at the beginning of the main function, but found a relatively insignificant improvement in the speedup (8.10x to 8.38x). Therefore, we concluded that this case obtains the speedup mostly from sparse CNN rather than parallelism.

w/o `cudaSetDevice(2);`

```

make rca56@Taipei:~/Sparse_CNN_C00$ make run
./CNNConvLayer
CPU time for executing a typical convolutional layer = 1331.14ms
GPU time for executing a typical convolutional layer = 164.37ms
Congratulations! You pass the check.
Speedup: 8.09844

```

w/ `cudaSetDevice(2);`

```

make rca56@Taipei:~/Sparse_CNN_C00$ make run
./CNNConvLayer
CPU time for executing a typical convolutional layer = 1334.05ms
GPU time for executing a typical convolutional layer = 159.195ms
Congratulations! You pass the check.
Speedup: 8.37995

```

### C. Show how you use NVVP to help you find and solve perf. Issues

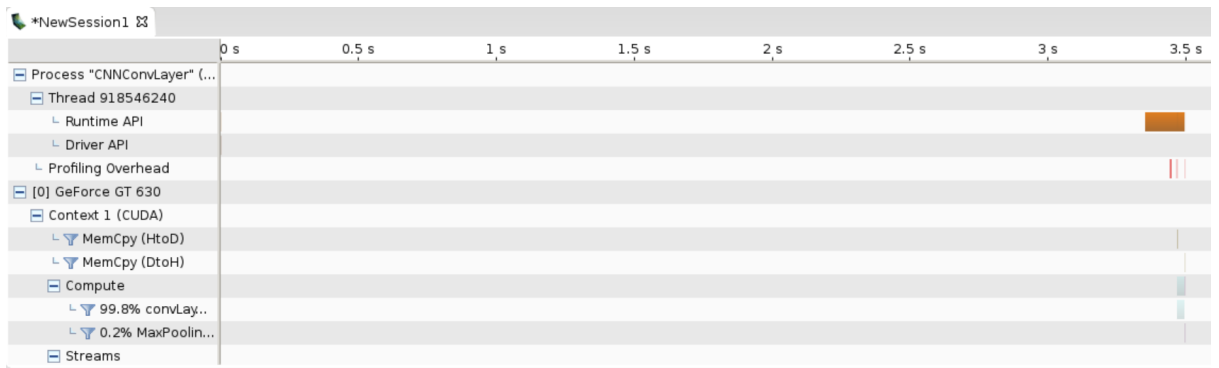
Checking NVVP of these two programs, both these two programs spend most time on **`cudaMalloc();`** and **`cudaDeviceSynchronize();`**. Thus, according to the Amdahl's Law, we improved **`cudaMalloc();`** by first access to GPU memory with **`cudaFree(0);`** before **`cudaMalloc();`**.

*(Explanation: Cuda is a lazy initialization, which means it won't give us its context until we first `cudaMalloc()` it. So, the way to reduce `cudaMalloc()` time is that we call `cudaFree(0)` first, and then it would have given us its context by the time we `cudaMalloc` it.)*

Before improvement, we can tell that **`cudaMalloc();`** and **`cudaDeviceSynchronize();`** spend most of time from following diagrams.

***Program 1: same as FP1***

*nvvp version*



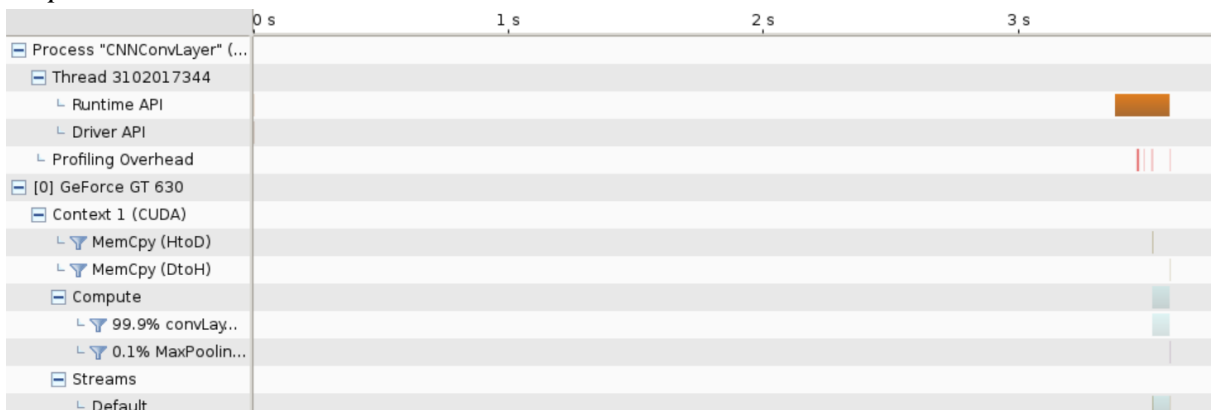
## nvprof version

```
[ca57@hsinchu CA2017FP-Part1]$ nvprof ./CNNConvLayer
==1021== NVPROF is profiling process 1021, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1654.07ms
GPU time for executing a typical convolutional layer = 303.853ms
Congratulations! You pass the check.
Speedup: 5.44366
==1021== Profiling application: ./CNNConvLayer
==1021== Profiling result:
Time(%)      Time       Calls      Avg       Min       Max      Name
96.91%    8.0077ms        1    8.0077ms    8.0077ms    8.0077ms    convLayerGPU(int*, int*, int*)
 2.87%    237.06us        2    118.53us   45.280us   191.78us    [CUDA memcpy HtoD]
 0.12%     9.5680us        1     9.5680us    9.5680us    9.5680us    PoolingGPU(int*, int*)
 0.11%     9.1200us        1     9.1200us    9.1200us    9.1200us    [CUDA memcpy DtoH]

==1021== API calls:
Time(%)      Time       Calls      Avg       Min       Max      Name
95.29%    295.31ms         4    73.828ms    7.5130us   294.62ms    cudaMalloc
 2.63%     8.1617ms         2     4.0809ms    16.715us    8.1450ms    cudaDeviceSynchronize
 1.41%     4.3651ms        364     11.991us    270ns     435.00us    cuDeviceGetAttribute
 0.43%     1.3421ms         4    335.52us   332.60us   341.27us    cuDeviceTotalMem
 0.12%     359.92us         4     89.979us    82.127us   101.61us    cuDeviceGetName
 0.09%     285.52us         3     95.172us    34.330us   196.63us    cudaMemcpy
 0.02%      55.361us         2     27.680us    21.409us   33.952us    cudaLaunch
 0.00%     14.639us         1     14.639us    14.639us   14.639us    cudaSetDevice
 0.00%      6.3390us        12         528ns     273ns     1.0170us    cuDeviceGet
 0.00%      5.6120us         4     1.4030us    557ns     3.8340us    cudaFree
 0.00%      5.0220us         3     1.6740us    396ns     3.6160us    cuDeviceGetCount
 0.00%      4.2660us         5         853ns     189ns     3.1030us    cudaSetupArgument
 0.00%      2.4140us         2     1.2070us    463ns     1.9510us    cudaConfigureCall
```

## Program 2: using COO format

### nvvp version





### nvprof version

```
ca56@Taipei:~/Sparse_CNN_C00$ nvprof ./CNNConvLayer
==9958== NVPROF is profiling process 9958, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1336.1ms
GPU time for executing a typical convolutional layer = 187.069ms
Congratulations! You pass the check.
Speedup: 7.14231
==9958== Profiling application: ./CNNConvLayer
==9958== Profiling result:
Time(%)    Time    Calls    Avg      Min      Max    Name
99.26%    68.171ms     1    68.171ms  68.171ms  68.171ms  convLayerGPUSparse(int*, int*, int*, int*, int*, int*)
0.66%    450.73us     5    90.146us  12.831us  112.80us  [CUDA memcpy HtoD]
0.06%    43.710us     1    43.710us  43.710us  43.710us  MaxPoolingGPU(int*, int*)
0.02%    14.591us     1    14.591us  14.591us  14.591us  [CUDA memcpy DtoH]

==9958== API calls:
Time(%)    Time    Calls    Avg      Min      Max    Name
63.04%    118.13ms     7    16.876ms  5.7980us  117.79ms  cudaMalloc
36.46%    68.313ms     2    34.156ms  46.481us  68.266ms  cudaDeviceSynchronize
0.29%    549.19us     6    91.531us  20.359us  171.90us  cudaMemcpy
0.13%    252.92us    91    2.7790us  146ns     109.04us  cuDeviceGetAttribute
0.02%    43.291us     2    21.645us  15.062us  28.229us  cudaLaunch
0.02%    35.392us     1    35.392us  35.392us  35.392us  cuDeviceTotalMem
0.02%    29.879us     1    29.879us  29.879us  29.879us  cuDeviceGetName
0.01%    16.961us     8    2.1200us  603ns     12.435us  cudaFree
0.00%    4.4910us     1    4.4910us  4.4910us  4.4910us  cudaSetDevice
0.00%    2.7470us     3      915ns    187ns     1.5170us  cuDeviceGetCount
0.00%    2.2560us     8      282ns    140ns      746ns  cudaSetupArgument
0.00%    1.8100us     2      905ns    617ns     1.1930us  cudaConfigureCall
0.00%    1.2240us     3      408ns    229ns      715ns  cuDeviceGet
```

Following are the results after accessing to cuda memory first with **cudaFree(0);**.

### Program 1: same as FPI

#### improved version

```
ca56@Taipei:~/Sparse_CNN_C00$ nvprof ./CNNConvLayer
==10391== NVPROF is profiling process 10391, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1329.38ms
GPU time for executing a typical convolutional layer = 28.043ms
Congratulations! You pass the check.
Speedup: 47.4049
==10391== Profiling application: ./CNNConvLayer
==10391== Profiling result:
Time(%)    Time    Calls    Avg      Min      Max    Name
98.10%    26.866ms     1    26.866ms  26.866ms  26.866ms  convLayerGPU(int*, int*, int*)
1.69%    462.83us     2    231.41us  86.908us  375.92us  [CUDA memcpy HtoD]
0.16%    43.327us     1    43.327us  43.327us  43.327us  MaxPoolingGPU(int*, int*)
0.05%    14.591us     1    14.591us  14.591us  14.591us  [CUDA memcpy DtoH]

==10391== API calls:
Time(%)    Time    Calls    Avg      Min      Max    Name
80.58%    118.10ms     9    13.122ms  639ns    118.09ms  cudaFree
18.59%    27.242ms     2    13.621ms  45.691us  27.196ms  cudaDeviceSynchronize
0.30%    432.59us     4    108.15us  6.4380us  210.58us  cudaMalloc
0.23%    341.34us    91    3.7500us  144ns     148.63us  cuDeviceGetAttribute
0.20%    292.31us     3    97.436us  36.989us  199.41us  cudaMemcpy
0.04%    51.523us     2    25.761us  14.642us  36.881us  cudaLaunch
0.03%    43.594us     1    43.594us  43.594us  43.594us  cuDeviceTotalMem
0.03%    38.932us     1    38.932us  38.932us  38.932us  cuDeviceGetName
0.00%    5.1520us     1    5.1520us  5.1520us  5.1520us  cudaSetDevice
0.00%    2.2120us     2    1.1060us  482ns     1.7300us  cudaConfigureCall
0.00%    2.1720us     3      724ns    161ns     1.4190us  cuDeviceGetCount
0.00%    1.8830us     5      376ns    140ns      739ns  cudaSetupArgument
0.00%    1.0190us     3      339ns    204ns     460ns  cuDeviceGet
```

The speedup has been boosted from about 5 to 47.

## Program 2: using COO format improved version

```
ca56@Taipei:~/Sparse_CNN_COO$ nvprof ./CNNConvLayer
==9497== NVPROF is profiling process 9497, command: ./CNNConvLayer
CPU time for executing a typical convolutional layer = 1332.35ms
GPU time for executing a typical convolutional layer = 69.459ms
Congratulations! You pass the check.
Speedup: 19.1818
==9497== Profiling application: ./CNNConvLayer
==9497== Profiling result:
Time(%)   Time      Calls      Avg      Min      Max      Name
99.26%   68.155ms      1   68.155ms   68.155ms   68.155ms   convLayerGPUSparse(int*, int*, int*, int*, int*, int*)
0.66%    450.60us      5   90.120us   12.831us   112.76us   [CUDA memcpy HtoD]
0.06%    44.446us      1   44.446us   44.446us   44.446us   MaxPoolingGPU(int*, int*)
0.02%    14.591us      1   14.591us   14.591us   14.591us   [CUDA memcpy DtoH]

==9497== API calls:
Time(%)   Time      Calls      Avg      Min      Max      Name
63.18%   119.92ms      9   13.325ms    596ns   119.91ms   cudaFree
35.98%    68.288ms      2   34.144ms   46.745us   68.241ms   cudaDeviceSynchronize
0.29%    556.99us      7   79.570us    6.6410us   196.65us   cudaMalloc
0.29%    545.46us      6   90.910us   21.565us   169.59us   cudaMemcpy
0.18%    340.55us     91    3.7420us    149ns   147.95us   cuDeviceGetAttribute
0.03%    49.320us      2   24.660us   15.196us   34.124us   cudaLaunch
0.02%    42.927us      1   42.927us   42.927us   42.927us   cuDeviceTotalMem
0.02%    37.902us      1   37.902us   37.902us   37.902us   cuDeviceGetName
0.00%    4.3460us      1    4.3460us    4.3460us    4.3460us   cudaSetDevice
0.00%    2.3460us      2    1.1730us    544ns    1.8020us   cudaConfigureCall
0.00%    2.2890us      8        286ns    142ns     614ns   cudaSetupArgument
0.00%    2.0870us      3        695ns    180ns    1.3300us   cuDeviceGetCount
0.00%    1.1870us      3        395ns    222ns     592ns   cuDeviceGet
```

The speedup has been boosted from about 8 to 19.

### D. Feedback of this part

In FP2, we had spent lots of time figuring out how to use COO format, and even given the frame paddings to meet the function judgement. Also, we cut less blocks of GPU this time so that the speedup of GPU is limited.

However, after improving the GPU memory access time, which is the largest part of GPU computing, we find out that the more blocks we divide, the quicker the programs can be. In the other words, using `cudaFree()`; to exclude memory access time improved the first one more than the second. And the best performance is the improved version of FP1, so we will hand in that one.